Posterior inference

Joachim Vandekerckhove

Posterior inference

- The posterior distribution gives a lot of information

Posterior inference

- The posterior distribution gives a lot of information
- Ideally, we could just show a graph of it and leave the interpretation to the reader

Posterior inference

- The posterior distribution gives a lot of information
- Ideally, we could just show a graph of it and leave the interpretation to the reader
  - But often the posterior will have many dimensions

Posterior inference

- The posterior distribution gives a lot of information
- Ideally, we could just show a graph of it and leave the interpretation to the reader
  - But often the posterior will have many dimensions
  - And also that seems lazy

Posterior inference

- The posterior distribution gives a lot of information
- Ideally, we could just show a graph of it and leave the interpretation to the reader
    - But often the posterior will have many dimensions
    - And also that seems lazy
- We need a way to describe the posterior distribution

Posterior inference

- The posterior distribution gives a lot of information
- Ideally, we could just show a graph of it and leave the interpretation to the reader
    - But often the posterior will have many dimensions
    - And also that seems lazy
- We need a way to describe the posterior distribution
    - Mean? SD? Skew? Kurtosis? Mass at or around a certain value? $p(.8 \leq P_R \leq .9 | \#R, \#W)$?

Determining these summary statistics analytically can be daunting and may in fact not be possible in general.

Posterior inference

Determining these summary statistics analytically can be daunting and may in fact not be possible in general.

Here, with the discrete domain of P_R:

$$p(.8 \leq P_R \leq .9 | \#R, \#W) =$$
$$p(P_R = .80 | \#R, \#W)$$
$$+ p(P_R = .85 | \#R, \#W)$$
$$+ p(P_R = .90 | \#R, \#W)$$

Posterior inference

- A more continuous case on $x = 2P_R - 1$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $x \in (0, 1)$ is a convenient way to say $P_R \in (0.5, 1)$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $x \in (0, 1)$ is a convenient way to say $P_R \in (0.5, 1)$
  - $p(x) = B(x|\alpha, \beta) = Kx^{\alpha-1}(1-x)^{\beta-1}$
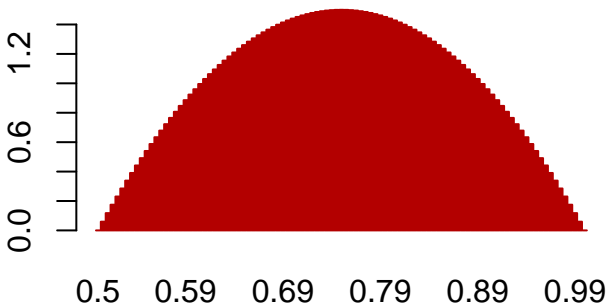
Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $x \in (0, 1)$ is a convenient way to say $P_R \in (0.5, 1)$
  - $p(x) = B(x|\alpha, \beta) = K x^{\alpha-1}(1-x)^{\beta-1}$
  - Beta distribution, say $\alpha = \beta = 2$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $x \in (0,1)$ is a convenient way to say $P_R \in (0.5, 1)$
  - $p(x) = B(x|\alpha, \beta) = Kx^{\alpha-1}(1-x)^{\beta-1}$
  - Beta distribution, say $\alpha = \beta = 2$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $x \in (0,1)$ is a convenient way to say $P_R \in (0.5, 1)$
  - $p(x) = B(x|\alpha, \beta) = Kx^{\alpha-1}(1-x)^{\beta-1}$
  - Beta distribution, say $\alpha = \beta = 2$

wine

- A more continuous case on $x = 2P_R - 1$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $p(x) = B(x|\alpha, \beta) = Kx^{\alpha-1}(1-x)^{\beta-1}$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $p(x) = B(x|\alpha, \beta) = Kx^{\alpha-1}(1-x)^{\beta-1}$
  - Beta distribution, say $\alpha = \beta = 2$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $p(x) = B(x|\alpha, \beta) = Kx^{\alpha-1}(1-x)^{\beta-1}$
  - Beta distribution, say $\alpha = \beta = 2$
  - Likelihood is the same: $p(\#R, \#W|P_R) = CP_R^{\#R}(1-P_R)^{\#W}$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $p(x) = B(x|\alpha, \beta) = Kx^{\alpha-1}(1-x)^{\beta-1}$
  - Beta distribution, say $\alpha = \beta = 2$
  - Likelihood is the same: $p(\#R, \#W|P_R) = CP_R^{\#R}(1-P_R)^{\#W}$
  - So the posterior must be:
    $p(P_R|\#R, \#W) \propto (2P_R - 1)^{\alpha-1}(1 - (2P_R - 1))^{\beta-1}$
    $\times P_R^{\#R}(1-P_R)^{\#W}$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $p(x) = B(x|\alpha, \beta) = K x^{\alpha-1} (1-x)^{\beta-1}$
  - Beta distribution, say $\alpha = \beta = 2$
  - Likelihood is the same: $p(\#R, \#W | P_R) = C P_R^{\#R} (1-P_R)^{\#W}$
  - So the posterior must be:
    $p(P_R | \#R, \#W) \propto (2P_R - 1)^{\alpha-1} (1 - (2P_R - 1))^{\beta-1}$
    $\times P_R^{\#R} (1-P_R)^{\#W}$

Posterior inference

- A more continuous case on $x = 2P_R - 1$
  - $p(x) = B(x|\alpha, \beta) = Kx^{\alpha-1}(1-x)^{\beta-1}$
  - Beta distribution, say $\alpha = \beta = 2$
  - Likelihood is the same: $p(\#R, \#W|P_R) = CP_R^{\#R}(1-P_R)^{\#W}$
  - So the posterior must be:
    $p(P_R|\#R, \#W) \propto (2P_R - 1)^{\alpha-1}(1 - (2P_R - 1))^{\beta-1}$
    $\times P_R^{\#R}(1-P_R)^{\#W}$

$\Rightarrow p(P_R|\#R, \#W) \propto (2P_R - 1)^{\alpha-1}(2 - 2P_R)^{\beta-1} P_R^{\#R}(1-P_R)^{\#W}$

. . . now what? What's the mean and variance of this?

Functional programming

Sometimes it is useful in R to turn a function into a variable to change it quickly

You can make a function "on the fly" inside a function or script file like this:

```
funcname <- function(n, x) { rep(x, n) }
```

So that prior and likelihood can be written like:

```
prior <- function(p) { dbeta(2*p - 1), 2, 2) }
likelihood <- function(p) { dbinom(5, 6, p) }
```

Prior and likelihood:

```
prior <- function(p) { dbeta(2*p - 1, 2, 2) }
likelihood <- function(p) { dbinom(5, 6, p) }
```
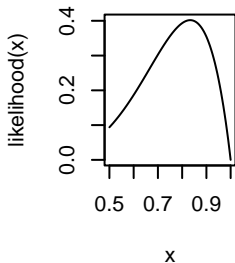
Given those, building the posterior is trivial:

```
post.usc  <- function(p) { prior(p) * likelihood(p) }
```

Exercise: implement this, plot the three functions

Functional programming

```
par(mfrow=c(1,3))
curve(prior, 0.5, 1)
curve(likelihood, 0.5, 1)
curve(post.usc, 0.5, 1)
```

Posterior inference

- These graphs still don't tell us the mean of the posterior (or any other useful statistic)

Posterior inference

- These graphs still don't tell us the mean of the posterior (or any other useful statistic)
- How do we determine the mean of an arbitrary, somewhat complicated function?

Posterior inference

- These graphs still don't tell us the mean of the posterior (or any other useful statistic)
- How do we determine the mean of an arbitrary, somewhat complicated function?
- As it turns out, drawing random samples from a distribution is an efficient way to do that

Posterior inference

- These graphs still don't tell us the mean of the posterior (or any other useful statistic)
- How do we determine the mean of an arbitrary, somewhat complicated function?
- As it turns out, drawing random samples from a distribution is an efficient way to do that
 - Methods for doing this are called Monte Carlo methods

Posterior inference

- These graphs still don't tell us the mean of the posterior (or any other useful statistic)
- How do we determine the mean of an arbitrary, somewhat complicated function?
- As it turns out, drawing random samples from a distribution is an efficient way to do that
    - Methods for doing this are called Monte Carlo methods
    - Math win: Monte Carlo methods don't need those hard-to-compute K and C scaling constants

Posterior sampling

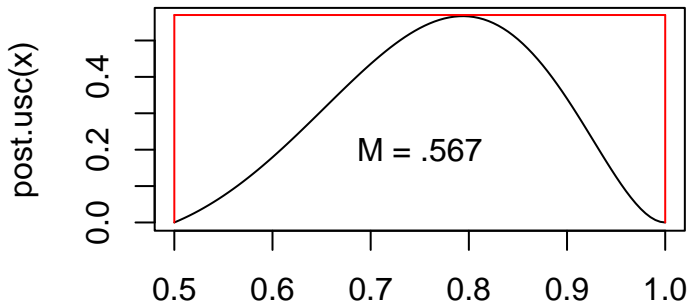One Monte Carlo method is the rejection sampler:

1) Draw a sample from some basic distribution $S(x|...)$
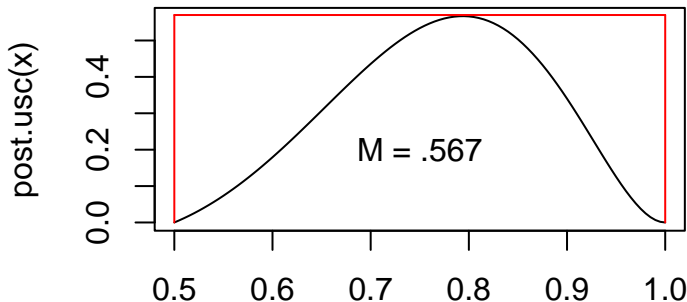
Posterior sampling

One Monte Carlo method is the rejection sampler:

1) Draw a sample from some basic distribution $S(x|...)$
2) Reject the sample with probability $q = \frac{fx}{M \times S(x)}$, where M is chosen so that this is always ≤ 1 (but ideally sometimes close to 1)

Posterior sampling

- Exercise: Implement a rejection algorithm

Posterior sampling

- Exercise: Implement a rejection algorithm

1. Sample P_R from $g(P_R) = U(0.5, 1.0)$ and u from $U(0, 1)$

Posterior sampling

- Exercise: Implement a rejection algorithm

1. Sample P_R from $g(P_R) = U(0.5, 1.0)$ and u from $U(0, 1)$
2. Check if $u < p(P_R|\#R, \#W)/M$ (or, better, if $uM < p(P_R|\#R, \#W)$)

Posterior sampling

- Exercise: Implement a rejection algorithm
1. Sample P_R from $g(P_R) = U(0.5, 1.0)$ and u from $U(0, 1)$
2. Check if $u < p(P_R | \#R, \#W)/M$ (or, better, if
   $uM < p(P_R | \#R, \#W))$
   - If true, accept x

Posterior sampling

- Exercise: Implement a rejection algorithm
1. Sample P_R from $g(P_R) = U(0.5, 1.0)$ and u from $U(0, 1)$
2. Check if $u < p(P_R|\#R, \#W)/M$ (or, better, if
   $uM < p(P_R|\#R, \#W)$)
   - If true, accept x
   - Otherwise, reject the value and draw a new sample

Posterior sampling

- Exercise: Implement a rejection algorithm
1. Sample P_R from $g(P_R) = U(0.5, 1.0)$ and u from $U(0,1)$
2. Check if $u < p(P_R | \#R, \#W)/M$ (or, better, if $uM < p(P_R | \#R, \#W)$)
   - If true, accept x
   - Otherwise, reject the value and draw a new sample
3. Repeat many times to get a few thousand samples

Posterior sampling

- Exercise: Implement a rejection algorithm
1. Sample P_R from $g(P_R) = U(0.5, 1.0)$ and u from $U(0,1)$
2. Check if $u < p(P_R|\#R, \#W)/M$ (or, better, if $uM < p(P_R|\#R, \#W)$)
   - If true, accept x
   - Otherwise, reject the value and draw a new sample
3. Repeat many times to get a few thousand samples
4. Make a histogram and compare the shapes of the distribution

Posterior sampling

```
N <- 10000
x <- vector(,N)
c <- 1
M <- .567
while(c <= N) {
  x[c] <- runif(1, 0.5, 1)
  u <- runif(1, 0, M)
  if (u < post.usc(x[c])) c <- c + 1;
}
hist(x, breaks=25,freq=FALSE)
K <- integrate(post.usc, 0.5, 1)$value
lines(domain, post.usc(domain)/K, lwd=3, col='red')
```

Histogram of x

- With a few thousand samples, the shape of the posterior is well approximated

- With a few thousand samples, the shape of the posterior is well approximated
  - Now we can compute the mean of that sample: 0.7708878

- With a few thousand samples, the shape of the posterior is well approximated
  - Now we can compute the mean of that sample: 0.7708878
  - . . . or the proportion of samples that are $> .85$: 0.2367

Posterior sampling

- With a few thousand samples, the shape of the posterior is well
  approximated
    - Now we can compute the mean of that sample: 0.7708878
    - ... or the proportion of samples that are $> .85$: 0.2367
    - ... or indeed any quality we fancy

Generative model representation

Generative model representation

- Basic unit of a Bayesian model is a distribution function

- Basic unit of a Bayesian model is a distribution function
- The posterior is defined through a generative model representation

- Basic unit of a Bayesian model is a distribution function
- The posterior is defined through a generative model representation
    - . . . which is basically a sequence of distributional assumptions

- Basic unit of a Bayesian model is a distribution function
- The posterior is defined through a generative model representation
    - . . . which is basically a sequence of distributional assumptions
    - In Bayesian statistics, a "model" is just a special kind of distribution function over parameters and data (and so with possibly very many dimensions)

Generative model representation

Let's define a really trivial model $\mathcal{M}_t$ in which we estimate the parameters μ and τ $(= 1/\sigma^2)$ of a normal distribution, applied to some data points d_j:

$$\mathcal{M}_t : \begin{cases} \forall j \in (1, \ldots, J) : \ d_j \sim N(\mu, \tau) \\ \mu \sim N(0, 0.1) \\ \tau \sim \Gamma(4, 0.01) \end{cases}$$

Let's define a really trivial model $\mathcal{M}_t$ in which we estimate the parameters μ and τ $(= 1/\sigma^2)$ of a normal distribution, applied to some data points d_j:

$$\mathcal{M}_t : \begin{cases} \forall j \in (1, \dots, J) : d_j \sim N(\mu, \tau) \\ \mu \sim N(0, 0.1) \\ \tau \sim \Gamma(4, 0.01) \end{cases}$$

Notice how every statement is a distributional assumption! (Either priors on parameters or likelihoods on data.)

Generative model representation

Let's define a really trivial model $\mathcal{M}_t$ in which we estimate the parameters μ and τ $(= 1/\sigma^2)$ of a normal distribution, applied to some data points d_j:

$$\mathcal{M}_t : \begin{cases} \forall j \in (1, \ldots, J) : \ d_j \sim N(\mu, \tau) \\ \mu \sim N(0, 0.1) \\ \tau \sim \Gamma(4, 0.01) \end{cases}$$

Notice how every statement is a distributional assumption! (Either priors on parameters or likelihoods on data.)

$$
\begin{aligned}
p(d_1, \ldots, d_J \mid \mathcal{M}_t) \ = \ & \left(\prod_{j=1}^{J} N(d_j \mid \mu, \tau) \right) \\
& \times \ N(\mu \mid 0, 0.1) \\
& \times \ \Gamma(\tau \mid 4, 0.01)
\end{aligned}
$$

JAGS code is (almost) perfect

$$\mathcal{M}_s : \begin{cases} \forall j \in (1, \ldots, J) : d_j \sim N(\mu, \tau) \\ \mu \sim N(0, 0.1) \\ \tau \sim \Gamma(4, 0.01) \end{cases}$$

The program needs to know the specifics of the model:

```
model {
    for (j in 1:J) {
        d[j] ~ dnorm(mu, tau)
    }
    mu ~ dnorm(0,0.1)
    tau ~ dgamma(4,0.01)
}
```

Summary and conclusion

The outcome of a Bayesian analysis is a probability distribution over parameters of interest

Summary and conclusion

The outcome of a Bayesian analysis is a probability distribution over parameters of interest

To communicate results, we needed a method to summarize arbitrary distributions

Summary and conclusion

The outcome of a Bayesian analysis is a probability distribution over parameters of interest

To communicate results, we needed a method to summarize arbitrary distributions

Monte Carlo methods provide such a method

Summary and conclusion

The outcome of a Bayesian analysis is a probability distribution over parameters of interest

To communicate results, we needed a method to summarize arbitrary distributions

Monte Carlo methods provide such a method

JAGS (and Stan) have a straightforward modeling language that lets users specify a model

Summary and conclusion

The outcome of a Bayesian analysis is a probability distribution over parameters of interest

To communicate results, we needed a method to summarize arbitrary distributions

Monte Carlo methods provide such a method

JAGS (and Stan) have a straightforward modeling language that lets users specify a model

They then draw samples from the posterior distributions, so we can calculate any summaries we like