

**If you're a Bayesian you can do everything that
God forbids (- Willem Heiser)**

Joachim Vandekerckhove

Unorthodox things you can do if you're a Bayesian

- Accumulate evidence for the absence of an effect

Unorthodox things you can do if you're a Bayesian

- Accumulate evidence for the absence of an effect
- Draw valid conclusions with almost no data

Unorthodox things you can do if you're a Bayesian

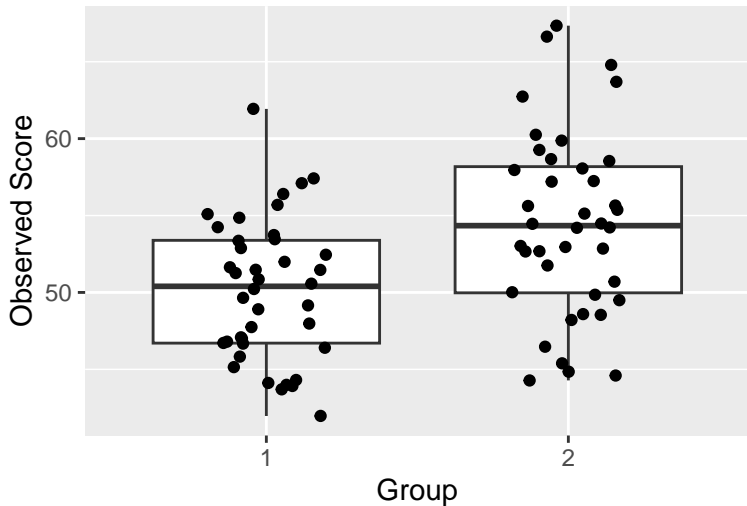
- Accumulate evidence for the absence of an effect
- Draw valid conclusions with almost no data
- Say things about means without knowing much about the basic units

Means without knowing much about the basic units

```
N <- 80
group <- rep(1:2, each = N/2) # Make two groups
# True score for each participant depends on group
true_score <-
  rnorm(N,
        mean = rep(c(50, 55), each = N/2),
        sd    = 2)
# Observations are noisy
observed_score <- true_score +
  rnorm(N, mean = 0, sd = 5)

data <- data.frame(participant = 1:N,
                   group = factor(group),
                   true_score, observed_score)
```

Observed Scores by Group



```
model_string <- "  
model {  
  for (i in 1:N) {  
    # Likelihood  
    observed_score[i] ~ dnorm(true_score[i], tau)  
    # Hierarchical model  
    true_score[i]      ~ dnorm(mu[group[i]], tau_true)  
  }  
  diff_mu ~ dnorm(0, 0.001)  
  mu[1]    ~ dunif(0, 100)  
  mu[2] <- mu[1] + diff_mu  
  tau      ~ dgamma(0.1, 0.1)  
  tau_true ~ dgamma(0.1, 0.1)  
}  
"
```

```
data_list <- list(  
  N = N,  
  observed_score = data$observed_score,  
  group = as.numeric(data$group)  
)  
  
# Initial values  
inits <- function() {  
  list(  
    tau = rgamma(1, 0.1, 0.1),  
    tau_true = rgamma(1, 0.1, 0.1),  
    true_score = runif(N, 0, 100)  
  )  
}
```



```
jags_model <- jags.model(textConnection(model_string),  
                          data      = data_list,  
                          inits     = inits,  
                          n.chains = 3 ,  
                          n.adapt  = 1000 )
```

```
## Compiling model graph  
##   Resolving undeclared variables  
##   Allocating nodes  
## Graph information:  
##   Observed stochastic nodes: 80  
##   Unobserved stochastic nodes: 84  
##   Total graph size: 250  
##  
## Initializing model
```

```

update(jags_model, n.iter = 1000)
samples <- coda.samples(jags_model,
                        variable.names = c("diff_mu",
                                           "true_score"),
                        n.iter = 5000)

```

```

# Summarize the results
summary_diff_mu <- summary(samples)$
  statistics["diff_mu", ]
print(summary_diff_mu)


```

##	Mean	SD	Naive SE	Time-series SE
##	4.314157098	1.203333239	0.009825175	0.058018522

```
# Extract true scores from the posterior samples
true_scores_posterior <- as.data.frame(as.matrix(samples)) %>%
  select(starts_with("true_score")) %>%
  apply(2, mean)

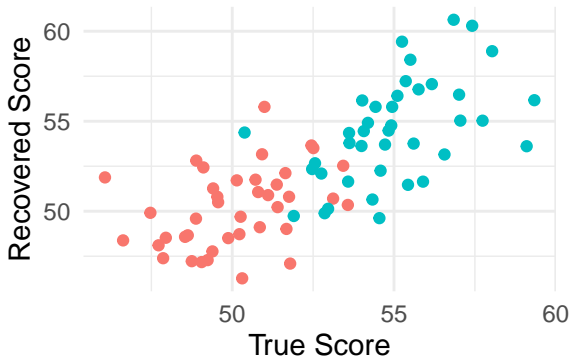
# Add the recovered true scores to the data frame
data$recovered_true_score <- true_scores_posterior

# Scatter plot of true scores vs recovered true scores
p2 <- ggplot(data, aes(x = true_score, y = recovered_true_score,
                        color = group)) +
  geom_point() +
  labs(x = "True Score", y = "Recovered Score") +
  theme_minimal() +
  theme(legend.position = "none")
```



Hierarchical recovery beats individual recovery

The difference between groups is well recovered as 4.3141571 ± 1.2033332 , even though the true scores within each group are poorly recovered.



Perform inference with no actual data

Suppose these were scores from a test (and suppose the two groups are an “on-track” group and an “advanced” group). One student from the advanced group missed class. What do we know about student $N + 1$?

```

model_string <- "
model {
  for (i in 1:N) {
    # Likelihood
    observed_score[i] ~ dnorm(true_score[i], tau)
    # Hierarchical model
    true_score[i] ~ dnorm(mu[group[i]], tau_true)
  }
  diff_mu ~ dnorm(0, 0.001)
  mu[1] ~ dunif(0, 100)
  mu[2] <- mu[1] + diff_mu
  tau ~ dgamma(0.1, 0.1)
  tau_true ~ dgamma(0.1, 0.1)

  true_score[N+1] ~ dnorm(mu[2], tau_true)
  observed_score[N+1] ~ dnorm(true_score[N+1], tau)
}

```

```
data_list <- list(  
  N = N,  
  observed_score = c(data$observed_score, NA),  
  group = as.numeric(data$group)  
)  
  
# Initial values  
inits <- function() {  
  list(  
    tau = rgamma(1, 0.1, 0.1),  
    tau_true = rgamma(1, 0.1, 0.1),  
    true_score = runif(N+1, 0, 100)  
  )  
}
```



```
jags_model <- jags.model(textConnection(model_string),  
                          data      = data_list,  
                          inits     = inits,  
                          n.chains = 3 ,  
                          n.adapt  = 1000 )
```

```
## Compiling model graph  
##   Resolving undeclared variables  
##   Allocating nodes  
## Graph information:  
##   Observed stochastic nodes: 80  
##   Unobserved stochastic nodes: 86  
##   Total graph size: 252  
##  
## Initializing model
```

```

update(jags_model, n.iter = 1000)
samples <- coda.samples(jags_model,
                        variable.names = c("diff_mu",
                                           "true_score",
                                           "observed_score"),
                        n.iter = 5000)

```

```

# Summarize the results
summary_new <- summary(samples)$
  statistics[c("true_score[81]", "observed_score[81]"),
            c("Mean", "SD")]
print(summary_new)

```

##	Mean	SD
## true_score[81]	54.41803	3.645873
## observed_score[81]	54.44141	5.458713

```
# Extract true scores from the posterior samples  
true_scores_posterior_new <- as.data.frame(as.matrix(samples)) %>%  
  select(starts_with("true_score[81]"))
```

